

# Machine Learning

## Regression. Classifier metrics

Aleksandr Petiushko

ML Research

August 12th, 2023



# Content

## ① Non-parametric Regression

# Content

- 1 Non-parametric Regression
- 2 Bias-Variance trade-off for k-NN Regression

# Content

- 1 Non-parametric Regression
- 2 Bias-Variance trade-off for k-NN Regression
- 3 Linear Regression

# Content

- ➊ Non-parametric Regression
- ➋ Bias-Variance trade-off for k-NN Regression
- ➌ Linear Regression
- ➍ ML and MAP Principles

# Content

- ➊ Non-parametric Regression
- ➋ Bias-Variance trade-off for k-NN Regression
- ➌ Linear Regression
- ➍ ML and MAP Principles
- ➎ Least Squares Method

# Content

- ➊ Non-parametric Regression
- ➋ Bias-Variance trade-off for k-NN Regression
- ➌ Linear Regression
- ➍ ML and MAP Principles
- ➎ Least Squares Method
- ➏ Ridge, LASSO and Elastic Net Regressions

# Content

- ➊ Non-parametric Regression
- ➋ Bias-Variance trade-off for k-NN Regression
- ➌ Linear Regression
- ➍ ML and MAP Principles
- ➎ Least Squares Method
- ➏ Ridge, LASSO and Elastic Net Regressions
- ➐ Regression Metrics



# Content

- ➊ Non-parametric Regression
- ➋ Bias-Variance trade-off for k-NN Regression
- ➌ Linear Regression
- ➍ ML and MAP Principles
- ➎ Least Squares Method
- ➏ Ridge, LASSO and Elastic Net Regressions
- ➐ Regression Metrics
- ➑ Classification Metrics and Confusion Matrix

# Content

- 1 Non-parametric Regression
- 2 Bias-Variance trade-off for k-NN Regression
- 3 Linear Regression
- 4 ML and MAP Principles
- 5 Least Squares Method
- 6 Ridge, LASSO and Elastic Net Regressions
- 7 Regression Metrics
- 8 Classification Metrics and Confusion Matrix
- 9 ROC and AUC

# Content

- 1 Non-parametric Regression
- 2 Bias-Variance trade-off for k-NN Regression
- 3 Linear Regression
- 4 ML and MAP Principles
- 5 Least Squares Method
- 6 Ridge, LASSO and Elastic Net Regressions
- 7 Regression Metrics
- 8 Classification Metrics and Confusion Matrix
- 9 ROC and AUC
- 10 Precision and Recall

# Content

- 1 Non-parametric Regression
- 2 Bias-Variance trade-off for k-NN Regression
- 3 Linear Regression
- 4 ML and MAP Principles
- 5 Least Squares Method
- 6 Ridge, LASSO and Elastic Net Regressions
- 7 Regression Metrics
- 8 Classification Metrics and Confusion Matrix
- 9 ROC and AUC
- 10 Precision and Recall
- 11 Multi-class case

# Non-parametric Regression

- The main disadvantage of parametric models is that it is necessary to have a parametric model to describe the dependency

# Non-parametric Regression

- The main disadvantage of parametric models is that it is necessary to have a parametric model to describe the dependency
- If it is impossible to select an adequate model, it makes sense to use non-parametric regression methods

# Non-parametric Regression

- The main disadvantage of parametric models is that it is necessary to have a parametric model to describe the dependency
- If it is impossible to select an adequate model, it makes sense to use non-parametric regression methods

## Assumption

Close objects correspond to close answers

# Non-parametric Regression

## The simplest model

We approximate the desired dependence by a constant in some neighborhood

## Nadaraya-Watson kernel regression<sup>1</sup>

If there are several objects from the training sample in the vicinity of the point, then it is reasonable to use the weighted average as a prediction of the algorithm

$$a(x) = \frac{\sum_i y_i \omega_i(x)}{\sum_i \omega_i(x)},$$

where  $\omega_i(x) = K_h(x, x_i)$ , a function  $K_h$  is called a **kernel** with smoothing window width  $h$ .

**Note:** we used for the homework  $\omega_i(x) = \frac{1}{k}$  for the k-NN method.

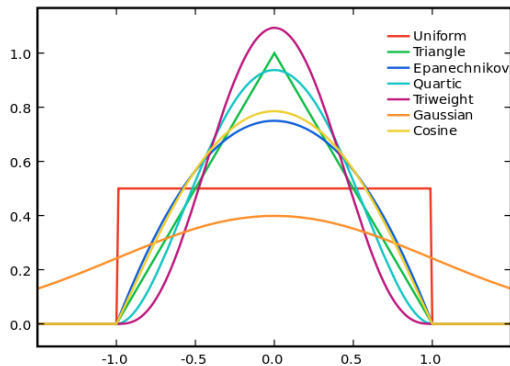
---

<sup>1</sup>[https://en.wikipedia.org/wiki/Kernel\\_regression](https://en.wikipedia.org/wiki/Kernel_regression)



# Examples of Kernels

- $K_h(x, x_i) = K\left(\frac{\|x - x_i\|}{h}\right)$
- Typical Examples: <sup>2</sup>



<sup>2</sup>[https://en.wikipedia.org/wiki/Kernel\\_\(statistics\)](https://en.wikipedia.org/wiki/Kernel_(statistics))

## Reminder: bias-variance tradeoff

### Definitions

Let  $y = y(x) = f(x) + \varepsilon$  be the target dependence, where  $f(x)$  is the deterministic function,  $\varepsilon \sim N(0, \sigma^2)$  and  $a(x)$  is the machine learning algorithm.

$$E(y - a)^2 = \sigma^2 + \text{variance}(a) + \text{bias}^2(f, a)$$

# Bias and Variance of k-NN Regression

## Bias

$$\text{bias}^2(f, a) = (E(f(x_0) - a(x_0)))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2$$

# Bias and Variance of k-NN Regression

## Bias

$$\text{bias}^2(f, a) = (E(f(x_0) - a(x_0)))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2$$

## Variance

$$\begin{aligned} \text{Variance}(a) &= D \left( \frac{1}{k} \sum_{i=1}^k y(x_{(i)}) \right) = \frac{1}{k^2} D \left( \sum_{i=1}^k y(x_{(i)}) \right) = \\ &= \frac{1}{k^2} D \left( \sum_{i=1}^k (f(x_{(i)}) + \varepsilon_i) \right) = \frac{1}{k^2} D \left( \sum_{i=1}^k f(x_{(i)}) \right) + \frac{1}{k^2} D \left( \sum_{i=1}^k \varepsilon_i \right) = \\ &= 0 + \frac{1}{k^2} k \sigma^2 = \frac{\sigma^2}{k} \end{aligned}$$

# Bias-Variance tradeoff of k-NN Regression

$$Error(x_0) = E(a(x_0) - f(x_0))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2 + \frac{\sigma^2}{k} + \sigma^2$$

## Bias-Variance tradeoff of k-NN Regression

$$Error(x_0) = E(a(x_0) - f(x_0))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2 + \frac{\sigma^2}{k} + \sigma^2$$

- Higher  $k$ , lower variance

# Bias-Variance tradeoff of k-NN Regression

$$Error(x_0) = E(a(x_0) - f(x_0))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2 + \frac{\sigma^2}{k} + \sigma^2$$

- Higher  $k$ , lower variance
- Higher  $k$ , higher bias

## Bias-Variance tradeoff of k-NN Regression

$$Error(x_0) = E(a(x_0) - f(x_0))^2 = \left( f(x_0) - \frac{1}{k} \sum_{i=1}^k f(x_{(i)}) \right)^2 + \frac{\sigma^2}{k} + \sigma^2$$

- Higher  $k$ , lower variance
- Higher  $k$ , higher bias

**Note:** Under “reasonable assumptions” the bias of the 1-NN estimator vanishes entirely as the size of the training set approaches infinity



# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model

# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model
- The method has a large number of variations to customize

# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model
- The method has a large number of variations to customize
  - ▶ Metric learning (e.g.,  $l_*$ -metric variations)

# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model
- The method has a large number of variations to customize
  - ▶ Metric learning (e.g.,  $l_*$ -metric variations)
  - ▶ Number of nearest neighbors  $k$

# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model
- The method has a large number of variations to customize
  - ▶ Metric learning (e.g.,  $l_*$ -metric variations)
  - ▶ Number of nearest neighbors  $k$
  - ▶ Weights in the weighted version of the method

# Conclusion

- The main advantage of non-parametric regression is the absence of assumptions about the form of the dependence model
- The method has a large number of variations to customize
  - ▶ Metric learning (e.g.,  $l_*$ -metric variations)
  - ▶ Number of nearest neighbors  $k$
  - ▶ Weights in the weighted version of the method
  - ▶ Smoothing window width

Time for questions



# Problem Statement: Linear Regression

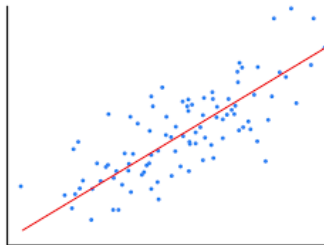
## Given

$$p(y_i|x_i) \sim w^T x_i + \varepsilon_i \sim N(w^T x_i, \sigma^2),$$

for  $i = 1.., \ell$ , where  $w \in \mathbf{R}^{n+1}$ ,  $\varepsilon_i \sim N(0, \sigma^2)$

## Task

Find  $w$





# Two kinds of parameter estimation

## Maximum Likelihood (**ML**) Principle

$$w_{ML} = \arg \max_w p(y|w, x)$$

# Two kinds of parameter estimation

## Maximum Likelihood (**ML**) Principle

$$w_{ML} = \arg \max_w p(y|w, x)$$

## Principle of Maximum A Posterior (**MAP**) Probability

$$w_{MAP} = \arg \max_w p(w|x, y)$$

# Maximum likelihood estimator

$$w_{ML} = \arg \max_w \sim p(y|w, x)$$

# Maximum likelihood estimator

$$w_{ML} = \arg \max_w \sim p(y|w, x)$$

$$w_{ML} = \arg \max_w \sim \prod_i p(y_i|w, x_i)$$

# Maximum likelihood estimator

$$w_{ML} = \arg \max_w \sim p(y|w, x)$$

$$w_{ML} = \arg \max_w \sim \prod_i p(y_i|w, x_i)$$

$$p(y_i|w, x_i) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)$$

# Maximum likelihood estimator

$$w_{ML} = \arg \max_w p(y|w, x)$$

$$w_{ML} = \arg \max_w \prod_i p(y_i|w, x_i)$$

$$p(y_i|w, x_i) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)$$

$$w_{ML} = \arg \max_w \prod_i \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right) = \arg \max_w \sum_i -\frac{(y_i - w^T x_i)^2}{2\sigma^2}$$

# Maximum likelihood estimator

$$w_{ML} = \arg \max_w p(y|w, x)$$

$$w_{ML} = \arg \max_w \prod_i p(y_i|w, x_i)$$

$$p(y_i|w, x_i) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)$$

$$w_{ML} = \arg \max_w \prod_i \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right) = \arg \max_w \sum_i -\frac{(y_i - w^T x_i)^2}{2\sigma^2}$$

$$w_{ML} = \arg \min_w \sum_i (y_i - w^T x_i)^2$$

# Least squares method

## Problem statement and assumptions

- $X = \mathbb{R}^n, Y = \mathbb{R}$



# Least squares method

## Problem statement and assumptions

- $X = \mathbb{R}^n$ ,  $Y = \mathbb{R}$
- $a(x) = f_w(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$ , where  $w = (w_0, w_1, \dots, w_n)^T \in \mathbb{R}^{n+1}$  — model parameters.

# Least squares method

## Problem statement and assumptions

- $X = \mathbb{R}^n$ ,  $Y = \mathbb{R}$
- $a(x) = f_w(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$ , where  $w = (w_0, w_1, \dots, w_n)^T \in \mathbb{R}^{n+1}$  — model parameters.
- It is convenient to write in vector form

$$a(x) = w^T \cdot x,$$

where  $x = (1, x^1, \dots, x^n)^T \in \mathbb{R}^{n+1}$ .

# Least squares method

## Problem statement and assumptions

- $X = \mathbb{R}^n$ ,  $Y = \mathbb{R}$
- $a(x) = f_w(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$ , where  $w = (w_0, w_1, \dots, w_n)^T \in \mathbb{R}^{n+1}$  — model parameters.
- It is convenient to write in vector form

$$a(x) = w^T \cdot x,$$

where  $x = (1, x^1, \dots, x^n)^T \in \mathbb{R}^{n+1}$ .

## Least squares method

- $L(w, X_{train}) = MSE(w, X_{train}) = \frac{1}{\ell} \sum_i (w^T \cdot x^{(i)} - y_i)^2$  — loss function

# Least squares method

## Problem statement and assumptions

- $X = \mathbb{R}^n$ ,  $Y = \mathbb{R}$
- $a(x) = f_w(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$ , where  $w = (w_0, w_1, \dots, w_n)^T \in \mathbb{R}^{n+1}$  — model parameters.
- It is convenient to write in vector form

$$a(x) = w^T \cdot x,$$

where  $x = (1, x^1, \dots, x^n)^T \in \mathbb{R}^{n+1}$ .

## Least squares method

- $L(w, X_{train}) = MSE(w, X_{train}) = \frac{1}{\ell} \sum_i (w^T \cdot x^{(i)} - y_i)^2$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

# Analytical solution

## Theorem

The solution to the problem  $\arg \min_w (\sum_{i=1}^{\ell} (w^T \cdot x_i - y_i)^2)$  is  $\hat{w} = (X^T X)^{-1} \cdot X^T \cdot y$ , where  $X_{i,j} = x_i^j$ ,  $y = (y_1, \dots, y_{\ell})$ .

# Analytical solution

## Theorem

The solution to the problem  $\arg \min_w (\sum_{i=1}^{\ell} (w^T \cdot x_i - y_i)^2)$  is  $\hat{w} = (X^T X)^{-1} \cdot X^T \cdot y$ , where  $X_{i,j} = x_i^j$ ,  $y = (y_1, \dots, y_{\ell})$ .

## Proof idea

Let's write the problem in vector form  $\|Xw - y\|^2 \rightarrow \min_w$ . The necessary condition for a minimum in matrix form is:

$$\frac{\partial}{\partial w} \|Xw - y\|^2 = 0$$

# Polynomial Regression

## Idea

It is possible to generate new features based on existing ones by applying non-linear functions

# Polynomial Regression

## Idea

It is possible to generate new features based on existing ones by applying non-linear functions

## Transformation examples

- Exponentiation
- Pairwise products
- Square root



# Pros and cons of linear regression

# Pros and cons of linear regression

## Advantages

- Simple algorithm, not computationally complex
- Linear regression is a well interpretable model
- Despite its simplicity, it can describe quite complex dependencies (for example, polynomials)

# Pros and cons of linear regression

## Advantages

- Simple algorithm, not computationally complex
- Linear regression is a well interpretable model
- Despite its simplicity, it can describe quite complex dependencies (for example, polynomials)

## Disadvantages

- The algorithm assumes that all features are numeric
- The algorithm assumes that the data is normally distributed, which is not always the case
- The algorithm is highly sensitive to outliers

Time for questions



# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w \sim p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

$$w_{MAP} = \arg \max_w \prod_i p(y_i|x_i, w)p(w)$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w \sim p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

$$w_{MAP} = \arg \max_w \sim \prod_i p(y_i|x_i, w)p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i \ln p(y_i|x_i, w) + \ln p(w)$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w \sim p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

$$w_{MAP} = \arg \max_w \sim \prod_i p(y_i|x_i, w)p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i \ln p(y_i|x_i, w) + \ln p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i -\frac{(y_i - w^T x_i)^2}{2\sigma^2} + \ell \ln p(w)$$



# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w \sim p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

$$w_{MAP} = \arg \max_w \sim \prod_i p(y_i|x_i, w)p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i \ln p(y_i|x_i, w) + \ln p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i -\frac{(y_i - w^T x_i)^2}{2\sigma^2} + \ell \ln p(w)$$

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \max_w \sim p(w|x_1, \dots, x_\ell, y_1, \dots, y_\ell)$$

$$w_{MAP} = \arg \max_w \sim \prod_i p(y_i|x_i, w)p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i \ln p(y_i|x_i, w) + \ln p(w)$$

$$w_{MAP} = \arg \max_w \sim \sum_i -\frac{(y_i - w^T x_i)^2}{2\sigma^2} + \ell \ln p(w)$$

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

An additional term appeared in the minimization problem, which depends only on the prior distribution on the weights  $w$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

Let's assume that  $p(w) \sim N(0, \tau^2)$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

Let's assume that  $p(w) \sim N(0, \tau^2)$

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \frac{\ell w^T w}{2\tau^2}$$

# Maximum A Posterior Probability Method

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \ell \ln p(w)$$

Let's assume that  $p(w) \sim N(0, \tau^2)$

$$w_{MAP} = \arg \min_w \sim \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \frac{\ell w^T w}{2\tau^2}$$

$$w_{MAP} = \arg \min_w \sim \frac{1}{\ell} \sum_i \frac{(y_i - w^T x_i)^2}{2\sigma^2} - \frac{1}{2\tau^2} \|w\|^2$$

# Ridge Regression

# Ridge Regression

## $L_2$ regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \frac{\alpha}{2} \sum_{i=0}^n w_i^2 = \frac{1}{\ell} \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \frac{\alpha}{2} \sum_{i=0}^n w_i^2$  — loss function



# Ridge Regression

## $L_2$ regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \frac{\alpha}{2} \sum_{i=0}^n w_i^2 = \frac{1}{\ell} \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \frac{\alpha}{2} \sum_{i=0}^n w_i^2$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

# Ridge Regression

## $L_2$ regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \frac{\alpha}{2} \sum_{i=0}^n w_i^2 = \frac{1}{\ell} \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \frac{\alpha}{2} \sum_{i=0}^n w_i^2$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

Note:  $L_2$  regularization and MAP with the normally distributed weights are the same!

# Ridge regression: solution

## Theorem

The solution of the problem  $\arg \min_w (\sum_{i=1}^{\ell} (w^T \cdot x^{(i)} - y_i)^2 + \alpha \sum_{i=0}^n w_i^2)$  is  $\hat{w} = (X^T X + \alpha I_{n+1})^{-1} \cdot X^T \cdot y$ , where  $X_{i,j} = x_i^j$ ,  $y = (y_1, \dots, y_{\ell})$ ,  $I_{n+1}$  is the identity matrix.

## Proof idea

Let's write the problem in vector form  $\|Xw - y\|^2 + \alpha \|w\|^2 \rightarrow \min_w$ . The necessary condition for a minimum in matrix form is:

$$\frac{\partial}{\partial w} ((Xw - y)^T \cdot (Xw - y) + \alpha w^T w) = 0$$

# Ridge Regression: Properties

- Regularization prevents model parameters from being too large
- In general, regularization provides better generalization ability
- More resistant to outliers
- A parameter has been added that can be configured using cross-validation

# Ridge Regression: Properties

- Regularization prevents model parameters from being too large
- In general, regularization provides better generalization ability
- More resistant to outliers
- A parameter has been added that can be configured using cross-validation

The probabilistic meaning of the  $\alpha$  parameter

$\alpha = \frac{1}{\tau^2}$ , where  $\tau$  is the standard deviation of the prior distribution on  $w$

# LASSO

## $L_1$ -regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \alpha \sum_{i=0}^n |w_i| = \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \alpha \sum_{i=0}^n |w_i|$  — loss function

# LASSO

## $L_1$ -regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \alpha \sum_{i=0}^n |w_i| = \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \alpha \sum_{i=0}^n |w_i|$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

# LASSO

## $L_1$ -regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + \alpha \sum_{i=0}^n |w_i| = \sum_i (w^T \cdot x^{(i)} - y_i)^2 + \alpha \sum_{i=0}^n |w_i|$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

## Properties

- This regularization provides feature selection
- No analytical solution



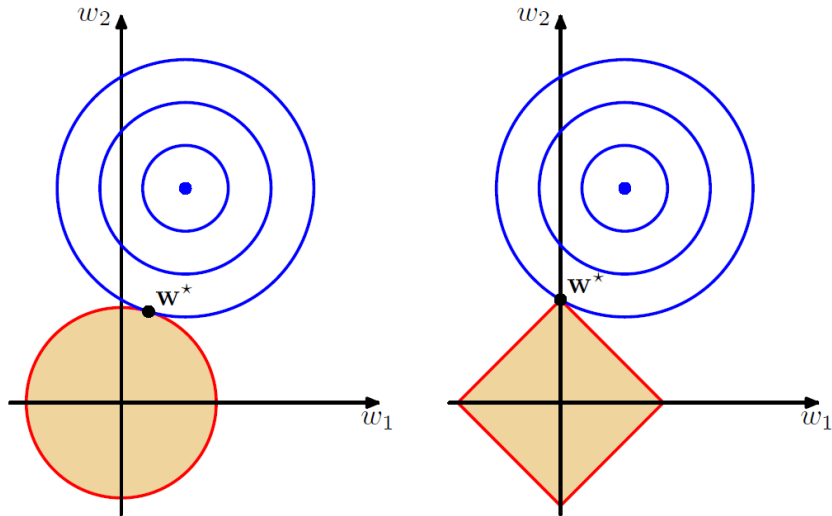
# Probabilistic interpretation of LASSO

## The probabilistic meaning of the $\alpha$ parameter

The parameter  $\alpha$  — is inversely proportional to the standard deviation of the prior distribution by  $w$ . In this case, this is the *Laplace* distribution

$$p(w) = \frac{1}{\tau} \exp\left(-\frac{\|w\|}{2\tau}\right)$$

## Intuition of feature selection under $L_1$ -regularization



## $L_1$ -regularization and $L_2$ -regularization

- $$L(w, X_{train}) = MSE(w, X_{train}) + r\alpha \sum_{i=0}^n |w_i| + (1-r)\frac{\alpha}{2} \sum_{i=0}^n w_i^2 =$$
$$\sum_i (w^T \cdot x^{(i)} - y_i)^2 + r\alpha \sum_{i=0}^n |w_i| + (1-r)\frac{\alpha}{2} \sum_{i=0}^n w_i^2 - \text{loss function}$$

# Elastic Net

## $L_1$ -regularization and $L_2$ -regularization

- $L(w, X_{train}) = MSE(w, X_{train}) + r\alpha \sum_{i=0}^n |w_i| + (1-r)\frac{\alpha}{2} \sum_{i=0}^n w_i^2 =$   
 $\sum_i (w^T \cdot x^{(i)} - y_i)^2 + r\alpha \sum_{i=0}^n |w_i| + (1-r)\frac{\alpha}{2} \sum_{i=0}^n w_i^2$  — loss function
- The task is to find  $\hat{w} = \arg \min_w (L(w, X_{train}))$

## Properties

- No analytical solution
- Combines the positive properties of Ridge regression and LASSO.

Time for questions



# Quality Metrics for the Regression Problem

## Motivation

- The formulation of a machine learning problem usually begins with the definition of a metric and fixing a test dataset on which this metric will be calculated

# Quality Metrics for the Regression Problem

## Motivation

- The formulation of a machine learning problem usually begins with the definition of a metric and fixing a test dataset on which this metric will be calculated
- An incorrectly chosen metric can make it difficult to use the machine learning model in real life and nullify the efforts of the team developing the machine learning algorithm

# Quality Metrics for the Regression Problem

## Motivation

- The formulation of a machine learning problem usually begins with the definition of a metric and fixing a test dataset on which this metric will be calculated
- An incorrectly chosen metric can make it difficult to use the machine learning model in real life and nullify the efforts of the team developing the machine learning algorithm
- As a rule, the customer does not think in terms of metrics and can only explain the problem he wants to solve in business language



# Quality Metrics for the Regression Problem

## Motivation

- The formulation of a machine learning problem usually begins with the definition of a metric and fixing a test dataset on which this metric will be calculated
- An incorrectly chosen metric can make it difficult to use the machine learning model in real life and nullify the efforts of the team developing the machine learning algorithm
- As a rule, the customer does not think in terms of metrics and can only explain the problem he wants to solve in business language
- Understanding the impact of the choice of a particular metric on the customer's business is the key to successful problem setting

# Quality Metrics for the Regression Problem

## Mean Square Error

$$MSE = \frac{1}{\ell} \sum_i (y_i - a(x_i))^2$$

# Quality Metrics for the Regression Problem

## Mean Square Error

$$MSE = \frac{1}{\ell} \sum_i (y_i - a(x_i))^2$$

## Root Mean Square Error

$$RMSE = \sqrt{\frac{1}{\ell} \sum_i (y_i - a(x_i))^2}$$

# Quality Metrics for the Regression Problem

## Mean Square Error

$$MSE = \frac{1}{\ell} \sum_i (y_i - a(x_i))^2$$

## Root Mean Square Error

$$RMSE = \sqrt{\frac{1}{\ell} \sum_i (y_i - a(x_i))^2}$$

## Mean Absolute Error

$$MAE = \frac{1}{\ell} \sum_i |y_i - a(x_i)|$$

# Quality Metrics for the Regression Problem

## Max Error

$$ME = \max(|y_i - a(x_i)|)$$

# Quality Metrics for the Regression Problem

## Max Error

$$ME = \max(|y_i - a(x_i)|)$$

## Mean Squared Logarithmic Error

$$MSLE = \frac{1}{\ell} \sum_i (\ln y_i - \ln a(x_i))^2$$

# Quality Metrics for the Regression Problem

## Max Error

$$ME = \max(|y_i - a(x_i)|)$$

## Mean Squared Logarithmic Error

$$MSLE = \frac{1}{\ell} \sum_i (\ln y_i - \ln a(x_i))^2$$

## $R^2$ score (also known as Coefficient of Determination)

$$R^2 = 1 - \frac{\sum_i (y_i - a(x_i))^2}{\sum_i (y_i - \bar{y})^2},$$

where  $\bar{y} = \frac{1}{\ell} \sum_i y_i$ .

# Conclusion

- Linear regression — simple, well-interpreted model, but not robust to outliers
- Has a clear probabilistic interpretation
- Regularization is a great way to deal with the overfitting and data noise



Time for questions



## Classification of binary classifier responses

- Training set  $X^m = \{(x_1, y_1), \dots, (x_m, y_m)\}$
- Classification problem into 2 classes:  $X \rightarrow Y, Y = \{+1, -1\}$
- Classification algorithm  $a(x) : X \rightarrow Y$
- The class labeled “+1” is called “**positive**”
- The class labeled “-1” is called “**negative**”

# Classification of binary classifier responses

- Training set  $X^m = \{(x_1, y_1), \dots, (x_m, y_m)\}$
- Classification problem into 2 classes:  $X \rightarrow Y, Y = \{+1, -1\}$
- Classification algorithm  $a(x) : X \rightarrow Y$
- The class labeled “+1” is called “**positive**”
- The class labeled “-1” is called “**negative**”

Таблица: Classification of responses




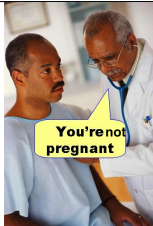
	Algorithm output	Correct answer
TP (True Positive)	$a(x_i) = +1$	$y_i = +1$
TN (True Negative)	$a(x_i) = -1$	$y_i = -1$
FP (False Positive)	$a(x_i) = +1$	$y_i = -1$
FN (False Negative)	$a(x_i) = -1$	$y_i = +1$

# Confusion Matrix

More clearly, these relationships can be depicted using **confusion matrix** (matrix of errors)

		Correct answer	
		$y = +1$	$y = -1$
Algorithm Output	$a(x) = +1$	True Positive	False Positive (Type 1 Error)
	$a(x) = -1$	False Negative (Type 2 Error)	True Negative

# Confusion Matrix

	$y = +1$	$y = -1$
$a(x) = +1$		
$a(x) = -1$		

# The simplest quality metric

- The simplest quality metric is the proportion of correct answers on a test (control sample)
- Common name: **Accuracy**

## Accuracy formula

$$Accuracy = \frac{1}{m} \sum_{i=1}^m [a(x_i) = y_i] = \frac{TP+TN}{TP+FP+TN+FN}$$

# The simplest quality metric

- The simplest quality metric is the proportion of correct answers on a test (control sample)
- Common name: **Accuracy**

## Accuracy formula

$$Accuracy = \frac{1}{m} \sum_{i=1}^m [a(x_i) = y_i] = \frac{TP+TN}{TP+FP+TN+FN}$$

## Disadvantages

- Ignores class imbalance
- The cost of an error on objects of different classes is not taken into account

# Metrics based on the positive response of the algorithm

Consider the metrics that are based on the calculation of the proportion of positive responses of the algorithm.

Proportion of *incorrect* positive classifications

Also known as False Positive Rate, or **FPR**.

$$FPR(a, X^m) = \frac{\sum_{i=1}^m [y_i = -1][a(x_i) = +1]}{\sum_{i=1}^m [y_i = -1]}$$



# Metrics based on the positive response of the algorithm

Consider the metrics that are based on the calculation of the proportion of positive responses of the algorithm.

## Proportion of *incorrect* positive classifications

Also known as False Positive Rate, or **FPR**.

$$FPR(a, X^m) = \frac{\sum_{i=1}^m [y_i = -1][a(x_i) = +1]}{\sum_{i=1}^m [y_i = -1]}$$

## Proportion of *correct* positive classifications

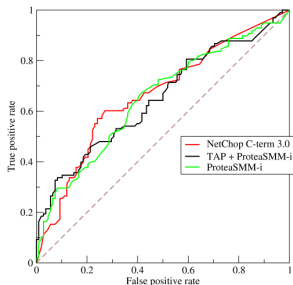
Also known as True Positive Rate, or **TPR**.

$$TPR(a, X^m) = \frac{\sum_{i=1}^m [y_i = +1][a(x_i) = +1]}{\sum_{i=1}^m [y_i = +1]}$$

**Note.** Notice the different denominators!

# Error Curve

Best known as Receiver Operating Characteristic (**ROC-curve**), in which we look at the trade-off between false alarm rate and correct response rate.



FPR is plotted along the X-axis, TPR is plotted along the Y-axis<sup>3</sup>.

**Note.** On this curve, miss rate (FN) is not taken into account in any way.

---

<sup>3</sup><https://wikipedia.org>

# Area under the ROC curve and types of ROC curves

## AUROC

The greater the value of the correct TPR prediction for each FPR error value, the better the classifier performs.

Thus, the area under the curve (Area Under Curve, AUC / AUROC) must be maximized.

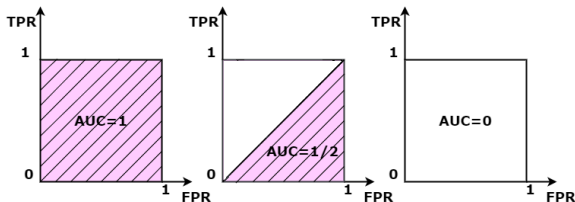
# Area under the ROC curve and types of ROC curves

## AUROC

The greater the value of the correct TPR prediction for each FPR error value, the better the classifier performs.

Thus, the area under the curve (Area Under Curve, AUC / AUROC) must be maximized.

ROC-curves for the best ( $AUC=1$ ), random ( $AUC=0.5$ ) and worst ( $AUC=0$ ) algorithm:



## The Task: build ROC, find AUROC

Suppose that the binary classification algorithm  $a(x_i)$  on the sample  $X^m$  decides to assign a class based on some scalar value  $g_\theta(x_i) \in \mathbb{R}$ , where  $\theta$  is the set of model parameters and  $g_\theta(x_i)$  is the discriminant function:

- Let's treat Positive response by a (varying) threshold  $t$ :  $g_\theta(x_i) \geq t$

### Task

- We want to build an ROC curve, i.e. find points  $\{(FPR_i, TPR_i)\}_{i=1}^m$
- Calculate area under curve - AUROC

## The Task: build ROC, find AUROC

Suppose that the binary classification algorithm  $a(x_i)$  on the sample  $X^m$  decides to assign a class based on some scalar value  $g_\theta(x_i) \in \mathbb{R}$ , where  $\theta$  is the set of model parameters and  $g_\theta(x_i)$  is the discriminant function:

- Let's treat Positive response by a (varying) threshold  $t$ :  $g_\theta(x_i) \geq t$

### Task

- We want to build an ROC curve, i.e. find points  $\{(FPR_i, TPR_i)\}_{i=1}^m$
- Calculate area under curve - AUROC

Let's count the number of correct answers of different types:

- $m_+ = \sum_{i=1}^m [y(x_i) = +1]$  (TPR denominator)
- $m_- = \sum_{i=1}^m [y(x_i) = -1]$  (FPR denominator);  $m = m_+ + m_-$

Let us order the training set  $X^m$  in descending order of the values  $g_\theta(x_i)$ .

Then the formula for  $AUROC = \frac{1}{m_-} \sum_{i=1}^m [y_i = -1] TPR_i$  (see below).

# Task solution

## Algorithm

We put the first point at the origin:  $(FPR_0, TPR_0) = (0, 0)$ ,  $AUROC = 0$ .

# Task solution

## Algorithm

We put the first point at the origin:  $(FPR_0, TPR_0) = (0, 0)$ ,  $AUROC = 0$ .

Loop over ordered selection  $i = 1 \dots m$

Threshold — the next value of the discriminant function  $t = g_\theta(x_i)$

If  $y_i = -1$ :

- $(FPR_i, TPR_i) = (FPR_{i-1} + \frac{1}{m_-}, TPR_{i-1})$  (move along the X-axis)
- $AUROC = AUROC + \frac{1}{m_-} TPR_i$



# Task solution

## Algorithm

We put the first point at the origin:  $(FPR_0, TPR_0) = (0, 0)$ ,  $AUROC = 0$ .

Loop over ordered selection  $i = 1 \dots m$

Threshold — the next value of the discriminant function  $t = g_\theta(x_i)$

If  $y_i = -1$ :

- $(FPR_i, TPR_i) = (FPR_{i-1} + \frac{1}{m_-}, TPR_{i-1})$  (move along the X-axis)
- $AUROC = AUROC + \frac{1}{m_-} TPR_i$

If  $y_i = +1$ :

- $(FPR_i, TPR_i) = (FPR_{i-1}, TPR_{i-1} + \frac{1}{m_+})$  (move along the Y-axis)

# Other Important Metrics 1

In information retrieval problems

- **Precision:**  $Precision = \frac{TP}{TP+FP}$  (percentage of relevant objects among those found)
- **Recall:**  $Recall = \frac{TP}{TP+FN}$  (percentage of found objects among relevant ones)

# Other Important Metrics 1

## In information retrieval problems

- **Precision:**  $Precision = \frac{TP}{TP+FP}$  (percentage of relevant objects among those found)
- **Recall:**  $Recall = \frac{TP}{TP+FN}$  (percentage of found objects among relevant ones)

## How to apply

- **Precision:** allows you to ensure that there are few false alarms; but it does not say anything about misses (the cost of a false alarm is high, and the price of a miss is low).
- **Recall:** allows you to ensure that there are few misses; but it does not say anything about false alarms (the price of a miss is high, and the price of a false alarm is low).

**Remark.** Often the task is to optimize one metric while fixing another.

## Other Important Metrics 2

In problems of medical diagnostics

- **Sensitivity:**  $Sensitivity = \frac{TP}{TP+FN}$  (percentage of correct positive diagnoses)
- **Specificity:**  $Specificity = \frac{TN}{TN+FP}$  (percentage of correct negative diagnoses)

## Other Important Metrics 2

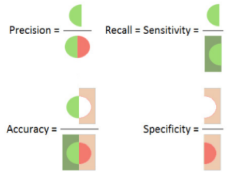
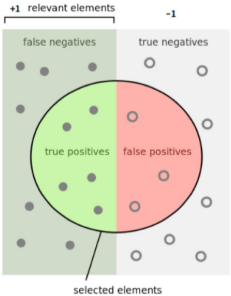
### In problems of medical diagnostics

- **Sensitivity:**  $Sensitivity = \frac{TP}{TP+FN}$  (percentage of correct positive diagnoses)
- **Specificity:**  $Specificity = \frac{TN}{TN+FP}$  (percentage of correct negative diagnoses)

### How to apply

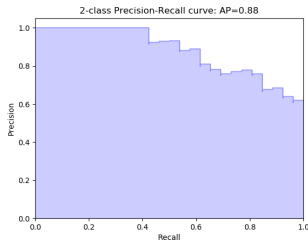
- **Sensitivity:** Maximize the number of true positive diagnoses, but ignore false diagnoses (treatment cost is low and skip cost is high).
- **Specificity:** Maximize the number of correct negative diagnoses, but don't take into account missed diagnoses (treatment cost is high and skip cost is low).

# Metrics illustration



# Aggregated Metrics over Precision-Recall

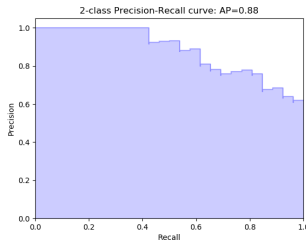
You can build a Precision-Recall (PR-curve) similar to the ROC-curve:



**Remark.** Note that in this case the curve is not necessarily monotonic!

# Aggregated Metrics over Precision-Recall

You can build a Precision-Recall (PR-curve) similar to the ROC-curve:



**Remark.** Note that in this case the curve is not necessarily monotonic!

## AUPRC

- Similar to AUROC, you can calculate the area under the PR curve - AUPRC
- Another name is Average Precision (with some assumptions on the integration method): the more, the better



# Multi-class classification

For each class  $c \in Y$ , denote by  $TP_c$ ,  $FP_c$ , and  $FN_c$  true positives, false positives, and false negatives. Then:

## Precision and recall with macro-averaging

- $Precision = \frac{\sum_c TP_c}{\sum_c (TP_c + FP_c)}$
- $Recall = \frac{\sum_c TP_c}{\sum_c (TP_c + FN_c)}$
- Insensitive to errors on small classes

# Multi-class classification

For each class  $c \in Y$ , denote by  $TP_c$ ,  $FP_c$ , and  $FN_c$  true positives, false positives, and false negatives. Then:

## Precision and recall with macro-averaging

- $Precision = \frac{\sum_c TP_c}{\sum_c (TP_c + FP_c)}$
- $Recall = \frac{\sum_c TP_c}{\sum_c (TP_c + FN_c)}$
- Insensitive to errors on small classes

## Precision and recall with micro-averaging

- $Precision = \frac{1}{|Y|} \sum_c \frac{TP_c}{TP_c + FP_c}$
- $Recall = \frac{1}{|Y|} \sum_c \frac{TP_c}{TP_c + FN_c}$
- Sensitive to errors on small classes

# Summary of classification quality metrics

- Precision and Recall are suitable for information retrieval tasks when the proportion of objects of the relevant class is small

# Summary of classification quality metrics

- Precision and Recall are suitable for information retrieval tasks when the proportion of objects of the relevant class is small
- Sensitivity and specificity are suitable for problems with unbalanced classes (as in medicine, for example)

# Summary of classification quality metrics

- Precision and Recall are suitable for information retrieval tasks when the proportion of objects of the relevant class is small
- Sensitivity and specificity are suitable for problems with unbalanced classes (as in medicine, for example)
- AUROC is suitable for quality assessment with a non-fixed error (miss rate) cost ratio

# Summary of classification quality metrics

- Precision and Recall are suitable for information retrieval tasks when the proportion of objects of the relevant class is small
- Sensitivity and specificity are suitable for problems with unbalanced classes (as in medicine, for example)
- AUROC is suitable for quality assessment with a non-fixed error (miss rate) cost ratio
- Another aggregated quality score - F-measure:  
$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$
  - ▶ This is the *harmonic mean* that goes to zero when at least one of the values goes to zero

Time for questions



# Thank you!